

Predicting Diabetes in the Pima Indians: An Investigation into Classification Strategies

Group 23: 490424010, 490390494

May 14, 2021

Contents

1	Introduction	3
1.1	Aim	3
1.2	Relevance	3
2	Data	4
2.1	Attribute Selection	4
3	Results and Discussion	5
3.1	Classifier Accuracy	5
3.2	DT Diagrams	5
3.3	Discussion	5
3.3.1	Comparison of Classifiers	5
3.3.2	Feature Selection	6
3.3.3	Decision Trees	6
3.3.4	Tree-based Classifiers	6
3.3.5	?Anything else that we consider important	7
4	Conclusion	8
5	Reflection	9
	Nomenclature	10
6	Appendix	11

List of Figures

1	The DT diagram of MyDT trained on the full discretised dataset.	15
2	The DT diagram of the Weka J48 algorithm <i>without</i> pruning and trained on the full discretised dataset.	16
3	The DT diagram of the Weka J48 algorithm <i>with</i> pruning and trained on the full discretised dataset.	16

List of Tables

1	A synopsis of the dataset's columns with those selected by CFS highlighted.	4
2	The 10-fold stratified cross validation accuracy in percentage (%) of each tested <i>numeric</i> classification algorithm using the dataset with and without CFS.	5

3	The 10-fold stratified cross validation accuracy in percentage (%) of each tested <i>nominal</i> classification algorithm using the dataset with and without CFS.	5
---	---	---

1 Introduction

1.1 Aim

The aim of this study is to investigate methods for predicting the future onset of diabetes mellitus (or simply diabetes), with relevance to females over 21 and of Pima Indian heritage.

1.2 Relevance

This study is important because the classifiers created have the potential to forewarn individuals of their risk to diabetes, and to be used as an easy clinical tool for early prevention. This is particularly important because, if left untreated, diabetes can lead to many serious long-term health implications such as cardiovascular disease, stroke, diabetic ketoacidosis and even death [1].

2 Data

The dataset used throughout this paper originates from the National Institute of Diabetes and Digestive and Kidney Diseases and was first used in a demonstration of the ADAP Learning Algorithm in 1988 [4]. It consists of 768 non-diabetic females aged at least 21 years old and of Pima Indian heritage. There are 9 columns per row, the first 8 of which are biometric measurement attributes whilst the final one is the class consisting of whether or not the individual will be diagnosed with diabetes. A description of each column in the dataset is shown in Table 1. To maintain consistency the dataset has been cleaned to remove any missing values.

Table 1: A synopsis of the dataset’s columns with those selected by CFS highlighted.

Description	Units
Number of times pregnant	n/a
Plasma glucose concentration at 2 hours in an oral glucose tolerance test	mg/dL
Diastolic blood pressure	mm Hg
Triceps skin fold thickness	mm
Serum insulin level	$\mu\text{U/mL}$
Body mass index (BMI)	kg/m^2
Diabetes pedigree function (likelihood of diabetes based on family history)	n/a
Age	years
Is diabetes diagnosed between 1 and 5 years after the above measurements are recorded?	n/a

2.1 Attribute Selection

The Correlation-based Feature Selection (CFS) method is a way of determining a representative set of attributes which are highly correlated with the class but uncorrelated with each other. This can improve the training of a classification model by removing features that are not predictive of the class.

Using the CFS algorithm [3] implemented in Weka 3.8.5 [2], the attributes that were selected are plasma glucose concentration, serum insulin level, BMI, diabetes pedigree function and age, and are additionally highlighted in Table 1.

3 Results and Discussion

3.1 Classifier Accuracy

The canonical Naïve Bayes (NB) and Decision Tree (DT) classification algorithms were implemented with tie decisions resulting in a ‘yes’ and are hereafter referred to as MyNB and MyDT respectively. 10-fold stratified cross validation was then performed on these algorithms and 12 other inbuilt Weka algorithms using the dataset described in section 2 after normalisation and discretisation for the numeric and nominal classification algorithms respectively.

Tables 2 and 3 present all the resulting accuracy figures for each tested classification algorithm, shown in percentage (%) to 4 d.p., using both the full dataset and the dataset after CFS, and coloured for ease of comparison.

Table 2: The 10-fold stratified cross validation accuracy in percentage (%) of each tested *numeric* classification algorithm using the dataset with and without CFS.

Numeric Data	ZeroR	1R	1NN	5NN	NB	MLP	SVM	MyNB
No feature selection	65.1042	70.8333	67.8385	74.4792	75.1302	75.3906	76.3021	75.2614
CFS	65.1042	70.8333	69.0104	74.4792	76.3021	75.7813	76.6927	76.0407

Table 3: The 10-fold stratified cross validation accuracy in percentage (%) of each tested *nominal* classification algorithm using the dataset with and without CFS.

Nominal Data	DT unpruned	DT pruned	MyDT	Bagg	Boost	RF
No feature selection	75.0000	75.3906	73.4484	74.8698	76.1719	73.1771
CFS	79.4271	79.4271	78.3869	78.5156	78.6458	78.9063

3.2 DT Diagrams

Decision trees were built on the full discretised dataset using three different algorithms: MyDT, and two DT classifiers from Weka (DT unpruned and DT pruned). The MyDT tree was built using the ID3 algorithm (without pruning), which recursively builds a tree based on maximum information gain. The two Weka variants were built using J48 (an implementation of the C4.5 algorithm) with default parameters, but differ in that one has been pruned in addition to the other [2]. The DT diagrams are displayed in Figures 1, 2 and 3 in section 6.

3.3 Discussion

3.3.1 Comparison of Classifiers

Overall, the accuracy of the 14 classifiers ranged roughly between 65% and 80% with a mean of $\sim 74.5\%$.

The best performing numeric classifier was the SVM, both with and without feature selection, where it achieved an accuracy of $\sim 76.7\%$ and $\sim 76.3\%$ respectively. Similar in performance were MyNB, NB and MLP, with accuracies roughly within 1% of the SVM. This small difference in accuracies ranging

from 75% to 77% is not necessarily indicative of algorithmic superiority but may be the effect of random noise in the testing dataset.

On the other hand, the worst performing numeric classifiers were ZeroR, 1R and 1NN, achieving accuracies between 65% and 71%. These simple algorithms are clearly not complex enough to capture patterns in the data, but are instead good points for comparison as to what is easily achievable (for example by predicting the majority class in ZeroR).

Within the nominal classifiers, the highest accuracy was $\sim 79.4\%$, and was obtained by both the pruned and unpruned DT using feature selection. Despite this, all of the nominal classifiers performed well using feature selection, with accuracies ranging roughly between 78% and 79.5%. Without feature selection, the best performing nominal classifier was Boost with an accuracy of $\sim 76.2\%$.

The worst performing nominal classifier was MyDT, with and without feature selection, where it achieved an accuracy of $\sim 78.3\%$ and $\sim 73.4\%$ respectively.

The 6 nominal classifiers clearly performed much better than the 8 numeric ones with a mean accuracy of $\sim 76.8\%$ compared to $\sim 72.8\%$. In addition, using CFS improved or equalled the performance of every classifier, with an average improvement in accuracy of $\sim 2.1\%$.

The implementations of MyNB and Weka's NB only differ in terms of their running time performance. In fact, the minimal differences in accuracies evident in Table 2 are most likely the result of different 10-fold data stratifications used in the cross validation accuracy calculations. On the other hand, the implementations of MyDT and Weka's two DTs differ profoundly. MyDT is built using the ID3 algorithm without pruning, whilst Weka uses J48 (an implementation of the 8th revision of the C4.5 algorithm [2]) which is very similar to ID3 but using the normalised information gain ratio as its splitting criterion. This resulted in Weka's two DTs performing better than MyDT with and without feature selection.

3.3.2 Feature Selection

glucose insulin bmi pedigree age

- improved all algos, some not as much (although this is covered in other places) - intuitive sense? maybe read literature and show these are important factors compared to others. also that these are mutually uncorrelated - can reference first couple splits in the decision and show that these are much same (eg glucose, bmi). therefore these are important factors, at least for information gain etc, so therefore correlated with class. also not so much correlated with each other, as the DTs would not see much gain in splitting on two similar variables A and B, as the first split would have removed much of correlation with class that A had, which would also remove correlation with class that B had, therefore no reason to split

3.3.3 Decision Trees

- similarity: glucose was used as first split for all trees, second split level is also similar - difference: much larger than equivalent unpruned, also less accurate suggesting overfitting - then segway into generic desc of pruning. how it works, how it leads to shorter tree and still has more accuracy

3.3.4 Tree-based Classifiers

how is this different from overall comparison of classifiers? basically just a comparison of nominal stuff dont wanna overlap too much i guess overall focuses on nominal vs numeric, and looks at best / worst overall

which DT method was used for dagg/boost/rf?

- boosting good even without CFS. try to speculate why. literature? is there a clear link between algos? boosting creates an iterative ensemble(?) of trees that focus on rows that we failed to predict, this is similar to having a number of uncorrelated features/trees QED? and then once CFS is used this advantage goes away - RF bad? if this uses very short trees we can blame this on inability to capture

complexity similar to numeric data. - read literature about DFS J48 to figure out why its much better than other algos / MyDT. prob just generically list the "improvements" over ID3 and just go therefore it performs better. - similar to bagging? bagging good bc reduces overfitting therefore good. if using small trees => still not able to fully capture complexity, therefore not as good as full J48 but better than RF. what tree does this use? if it uses full ID3 trees then this doesn't hold since its worse than them.

3.3.5 ?Anything else that we consider important

Nominal better? Weird?? Discussion point? or is the data being predicted here actually different? if not its probs just overfitting (to noise) or something when given more DOF(?) and thats something to mention.

could also talk about why J48 DT is the best. again, look into specifics of J48 and try to justify that it had all advantages of DT without disadvantages (+ advantages that other algos had).

if ur reading this then im already dead. jks im super busy until like 7pm today so ill turn these into actual paras when i get back.

4 Conclusion

conclusion

5 Reflection

References

- [1] AE, K., GE, U., JM, M., AND JN, F. Hyperglycemic crises in adult patients with diabetes. *Diabetes Care* 32, 7 (Jul 2009), 1335–1343.
- [2] FRANK, E., HALL, M. A., AND WITTEN, I. H. *Online Appendix for "Data Mining: Practical Machine Learning Tools and Techniques"*, 4 ed. Morgan Kaufmann, 2016.
- [3] HALL, M. A. *Correlation-based Feature Selection for Machine Learning*. PhD thesis, The University of Waikato, Apr 1999.
- [4] SMITH, J., EVERHART, J., DICKSON, W., KNOWLER, W., AND JOHANNES, R. Using the adap learning algorithm to forecast the onset of diabetes mellitus. *Proceedings - Annual Symposium on Computer Applications in Medical Care* 10 (Nov 1988).

Nomenclature

$\mu\text{U/mL}$ Micro enzyme units per millilitre

1R One Rule

Bagg Bagging

Boost Boosting

CFS Correlation-based feature selection

d.p. decimal points

DT Decision Tree

kg/m^2 Weight in kilograms per height in metres squared

kNN k-Nearest Neighbours

mg/dL Milligrams per decilitre

MLP Multilayer Perceptron

mm Millimetres

mm Hg Millimetres of mercury

n/a Not applicable

NB Naïve Bayes

RF Random Forest

SVM Support Vector Machines

ZeroR Classifier that always predicts the majority class

6 Appendix

glucose = very high
| insulin = high
| | bmi = high
| | | npreg = high
| | | | pedigree = high: yes (16.0/0.0)
| | | | pedigree = low
| | | | blood = high: yes (12.0/3.0)
| | | | blood = low: yes (3.0/1.0)
| | | npreg = low
| | | | age = high
| | | | | pedigree = high
| | | | | triceps = high
| | | | | | blood = high: yes (10.0/5.0)
| | | | | | blood = low: yes (1.0/1.0)
| | | | | | triceps = low: yes (1.0/0.0)
| | | | | pedigree = low
| | | | | blood = high
| | | | | | triceps = high: yes (14.0/2.0)
| | | | | | triceps = low: yes (3.0/0.0)
| | | | | blood = low
| | | | | | triceps = high: yes (3.0/0.0)
| | | | | | triceps = low: yes (1.0/1.0)
| | | | age = low
| | | | | pedigree = high: yes (12.0/0.0)
| | | | | pedigree = low
| | | | | triceps = high
| | | | | | blood = high: yes (7.0/2.0)
| | | | | | blood = low: yes (3.0/0.0)
| | | | | triceps = low
| | | | | | blood = high: yes (1.0/0.0)
| | | | | | blood = low: no (1.0/0.0)
| | | bmi = low
| | | age = high
| | | | triceps = high
| | | | | npreg = high
| | | | | pedigree = high
| | | | | | blood = high: yes (1.0/1.0)
| | | | | | blood = low: yes (1.0/0.0)
| | | | | pedigree = low: yes (2.0/0.0)
| | | | | npreg = low
| | | | | | blood = high: yes (1.0/1.0)
| | | | | | blood = low: yes (1.0/1.0)
| | | | | triceps = low: yes (3.0/0.0)
| | | age = low
| | | | blood = high
| | | | | triceps = high: no (1.0/0.0)
| | | | | triceps = low: yes (1.0/1.0)
| | | | blood = low: no (1.0/0.0)
| insulin = low
| | pedigree = high: yes (1.0/0.0)
| | pedigree = low: no (2.0/0.0)

```

glucose = high
| bmi = high
| | age = high
| | | pedigree = high
| | | | blood = high
| | | | | npreg = high: yes (12.0/2.0)
| | | | | npreg = low
| | | | | | triceps = high
| | | | | | | insulin = high: yes (9.0/4.0)
| | | | | | | insulin = low: yes (1.0/0.0)
| | | | | | | triceps = low: yes (1.0/0.0)
| | | | | blood = low
| | | | | | npreg = high: no (2.0/0.0)
| | | | | | npreg = low: yes (1.0/1.0)
| | | | | pedigree = low
| | | | | | insulin = high
| | | | | | triceps = high
| | | | | | | blood = high
| | | | | | | | npreg = high: yes (9.0/6.0)
| | | | | | | | npreg = low: no (12.0/11.0)
| | | | | | | | blood = low
| | | | | | | | npreg = high: yes (3.0/0.0)
| | | | | | | | npreg = low: yes (2.0/1.0)
| | | | | | | | triceps = low
| | | | | | | | npreg = high: no (1.0/0.0)
| | | | | | | | npreg = low
| | | | | | | | | blood = high: yes (1.0/1.0)
| | | | | | | | | blood = low: no (1.0/0.0)
| | | | | | | | insulin = low: yes (1.0/0.0)
| | | | | age = low
| | | | | | triceps = high
| | | | | | | pedigree = high
| | | | | | | | blood = high: yes (6.0/5.0)
| | | | | | | | blood = low: no (3.0/2.0)
| | | | | | | | pedigree = low
| | | | | | | | | blood = high: no (12.0/8.0)
| | | | | | | | | blood = low: yes (4.0/3.0)
| | | | | | | | triceps = low
| | | | | | | | | blood = high: no (3.0/0.0)
| | | | | | | | | blood = low
| | | | | | | | | insulin = high
| | | | | | | | | | pedigree = high: yes (1.0/1.0)
| | | | | | | | | | pedigree = low: no (1.0/0.0)
| | | | | | | | | insulin = low: no (1.0/0.0)
| | | | | bmi = low
| | | | | | triceps = high
| | | | | | | insulin = high
| | | | | | | | pedigree = high: no (5.0/0.0)
| | | | | | | | pedigree = low
| | | | | | | | | age = high
| | | | | | | | | | blood = high: no (5.0/1.0)
| | | | | | | | | | blood = low: no (2.0/0.0)
| | | | | | | | | age = low

```

```

| | | | | blood = high: no (1.0/0.0)
| | | | | blood = low: yes (2.0/2.0)
| | | insulin = low
| | | | pedigree = high: yes (1.0/0.0)
| | | | pedigree = low: no (1.0/0.0)
| | | triceps = low: no (9.0/0.0)
glucose = medium
| age = high
| | bmi = high
| | | pedigree = high
| | | | npreg = high: yes (13.0/0.0)
| | | | npreg = low
| | | | | triceps = high
| | | | | blood = high: yes (9.0/7.0)
| | | | | blood = low: yes (3.0/3.0)
| | | | | triceps = low: yes (2.0/0.0)
| | | | pedigree = low
| | | | | insulin = high
| | | | | blood = high
| | | | | | npreg = high: no (14.0/12.0)
| | | | | | npreg = low
| | | | | | triceps = high: no (18.0/11.0)
| | | | | | triceps = low: yes (1.0/1.0)
| | | | | blood = low
| | | | | | triceps = high
| | | | | | npreg = high: yes (3.0/3.0)
| | | | | | npreg = low: yes (5.0/4.0)
| | | | | | triceps = low: no (2.0/1.0)
| | | | | insulin = low: no (5.0/0.0)
| | | bmi = low
| | | | blood = high
| | | | | npreg = high: no (13.0/0.0)
| | | | | npreg = low
| | | | | | pedigree = high: no (3.0/0.0)
| | | | | | pedigree = low
| | | | | | triceps = high: no (2.0/1.0)
| | | | | | triceps = low: no (2.0/0.0)
| | | | blood = low
| | | | | npreg = high: yes (1.0/0.0)
| | | | | npreg = low
| | | | | | triceps = high: no (5.0/0.0)
| | | | | | triceps = low: no (2.0/1.0)
| age = low
| | bmi = high
| | | triceps = high
| | | | npreg = high: yes (1.0/1.0)
| | | | npreg = low
| | | | | pedigree = high
| | | | | | blood = high
| | | | | | insulin = high: no (12.0/2.0)
| | | | | | insulin = low: no (3.0/0.0)
| | | | | blood = low
| | | | | | insulin = high: yes (3.0/3.0)

```

```

| | | | | insulin = low: yes (1.0/0.0)
| | | | | pedigree = low
| | | | | blood = high
| | | | | insulin = high: no (20.0/5.0)
| | | | | insulin = low: no (3.0/0.0)
| | | | | blood = low
| | | | | insulin = high: no (18.0/2.0)
| | | | | insulin = low: no (5.0/1.0)
| | | triceps = low
| | | | pedigree = high
| | | | blood = high: no (4.0/0.0)
| | | | blood = low
| | | | | insulin = high: no (3.0/1.0)
| | | | | insulin = low: no (2.0/0.0)
| | | | pedigree = low: no (14.0/0.0)
| | | bmi = low
| | | | pedigree = high
| | | | insulin = high: no (5.0/0.0)
| | | | insulin = low
| | | | | blood = high: no (1.0/0.0)
| | | | | blood = low: yes (1.0/1.0)
| | | | pedigree = low: no (34.0/0.0)
glucose = low
| | | bmi = high
| | | | insulin = high
| | | | | age = high
| | | | | pedigree = high
| | | | | blood = high
| | | | | npreg = high
| | | | | | triceps = high: yes (2.0/1.0)
| | | | | | triceps = low: no (1.0/0.0)
| | | | | npreg = low
| | | | | | triceps = high: no (1.0/0.0)
| | | | | | triceps = low: yes (1.0/0.0)
| | | | | blood = low: yes (1.0/0.0)
| | | | | pedigree = low
| | | | | triceps = high
| | | | | npreg = high
| | | | | | blood = high: no (8.0/1.0)
| | | | | | blood = low: no (1.0/0.0)
| | | | | npreg = low
| | | | | | blood = high: no (11.0/2.0)
| | | | | | blood = low: no (3.0/1.0)
| | | | | triceps = low: no (1.0/0.0)
| | | | age = low
| | | | | blood = high: no (18.0/0.0)
| | | | | blood = low
| | | | | triceps = high
| | | | | | pedigree = high: no (5.0/1.0)
| | | | | | pedigree = low: no (9.0/3.0)
| | | | | triceps = low: no (7.0/0.0)
| | | | insulin = low
| | | | | blood = high

```

```

| | | | age = high: no (12.0/0.0)
| | | | age = low
| | | | | triceps = high
| | | | | | pedigree = high: yes (1.0/0.0)
| | | | | | pedigree = low: no (5.0/1.0)
| | | | | | triceps = low: no (6.0/0.0)
| | | | blood = low: no (23.0/0.0)
| | | bmi = low: no (66.0/0.0)

```

Figure 1: The DT diagram of MyDT trained on the full discretised dataset.

```

glucose = high
| bmi = high
| | triceps = high
| | | npreg = low
| | | | pedigree = high
| | | | | age = high: yes (16.0/5.0)
| | | | | age = low
| | | | | | blood = high: yes (11.0/5.0)
| | | | | | blood = low: no (5.0/2.0)
| | | | | | pedigree = low
| | | | | | blood = high: no (43.0/19.0)
| | | | | | blood = low: yes (10.0/4.0)
| | | | npreg = high
| | | | | blood = high: yes (29.0/8.0)
| | | | | blood = low
| | | | | | pedigree = high: no (2.0)
| | | | | | pedigree = low: yes (3.0)
| | | | | triceps = low: no (13.0/4.0)
| | | | bmi = low: no (29.0/4.0)
glucose = low
| bmi = high
| | insulin = high
| | | age = high
| | | | pedigree = high: yes (7.0/3.0)
| | | | pedigree = low: no (28.0/4.0)
| | | | age = low: no (43.0/4.0)
| | | insulin = low: no (48.0/2.0)
| | | bmi = low: no (66.0)
glucose = very high
| insulin = high
| | bmi = high: yes (103.0/16.0)
| | | bmi = low
| | | | age = high: yes (12.0/3.0)
| | | | age = low: no (4.0/1.0)
| | | insulin = low: no (3.0/1.0)
glucose = medium
| age = high
| | insulin = high
| | | bmi = high
| | | | pedigree = high: yes (37.0/10.0)
| | | | pedigree = low
| | | | | blood = high: no (57.0/24.0)

```

```

| | | | blood = low
| | | | | triceps = high: yes (15.0/7.0)
| | | | | triceps = low: no (3.0/1.0)
| | | bmi = low: no (27.0/3.0)
| | insulin = low: no (8.0)
| age = low
| | bmi = high
| | | npreg = low
| | | | triceps = high
| | | | | pedigree = high
| | | | | blood = high: no (17.0/2.0)
| | | | | blood = low: yes (7.0/3.0)
| | | | | pedigree = low: no (54.0/8.0)
| | | | triceps = low: no (24.0/1.0)
| | | npreg = high: yes (2.0/1.0)
| | bmi = low: no (42.0/1.0)

```

Figure 2: The DT diagram of the Weka J48 algorithm *without* pruning and trained on the full discretised dataset.

```

glucose = high
| bmi = high
| | triceps = high: yes (119.0/51.0)
| | triceps = low: no (13.0/4.0)
| bmi = low: no (29.0/4.0)
glucose = low: no (192.0/14.0)
glucose = very high: yes (122.0/24.0)
glucose = medium
| age = high
| | bmi = high
| | | pedigree = high: yes (37.0/10.0)
| | | pedigree = low: no (80.0/33.0)
| | bmi = low: no (30.0/3.0)
| age = low: no (146.0/17.0)

```

Figure 3: The DT diagram of the Weka J48 algorithm *with* pruning and trained on the full discretised dataset.